

cm010 Exercises: Tibble Joins

Requirements

You will need Joey's `singer` R package for this exercise. And to install that, you'll need to install `devtools`. Running this code in your console should do the trick:

```
install.packages("devtools")
devtools::install_github("JoeyBernhardt/singer")
```

If you can't install the `singer` package, we've put the singer data on the `STAT545-UBC/Classroom` repo, and you can load those in instead by running the following lines of code:

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.2.1
## v ggplot2 3.2.1      v purrr  0.3.2
## v tibble  2.1.3      v dplyr  0.8.3
## v tidyr   1.0.0      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0
## Warning: package 'ggplot2' was built under R version 3.5.2
## Warning: package 'tibble' was built under R version 3.5.2
## Warning: package 'tidyr' was built under R version 3.5.2
## Warning: package 'purrr' was built under R version 3.5.2
## Warning: package 'dplyr' was built under R version 3.5.2
## Warning: package 'stringr' was built under R version 3.5.2
## Warning: package 'forcats' was built under R version 3.5.2
## -- Conflicts ----- tidyverse_conflicts()
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
songs <- read_csv("https://raw.githubusercontent.com/STAT545-UBC/Classroom/master/data/singer/songs.csv")

## Parsed with column specification:
## cols(
##   title = col_character(),
##   artist_name = col_character(),
##   year = col_double()
## )
locations <- read_csv("https://raw.githubusercontent.com/STAT545-UBC/Classroom/master/data/singer/locations.csv")

## Parsed with column specification:
## cols(
##   artist_name = col_character(),
##   city = col_character(),
##   release = col_character(),
##   title = col_character()
## )
```

Load required packages:

Exercise 1: singer

The package `singer` comes with two smallish data frames about songs. Let's take a look at them (after minor modifications by renaming and shuffling):

```
(time <- as_tibble(songs) %>%  
  rename(song = title))
```

```
## # A tibble: 22 x 3  
##   song                artist_name    year  
##   <chr>              <chr>      <dbl>  
## 1 Corduroy          Pearl Jam    1994  
## 2 Grievance         Pearl Jam    2000  
## 3 Stupidmop         Pearl Jam    1994  
## 4 Present Tense     Pearl Jam    1996  
## 5 MFC               Pearl Jam    1998  
## 6 Lukin             Pearl Jam    1996  
## 7 It's Lulu         The Boo Radleys 1995  
## 8 Sparrow           The Boo Radleys 1992  
## 9 Martin_ Doom! It's Seven O'Clock The Boo Radleys 1995  
## 10 Leaves And Sand  The Boo Radleys 1993  
## # ... with 12 more rows
```

```
(album <- as_tibble(locations) %>%  
  select(title, everything()) %>%  
  rename(album = release,  
         song = title))
```

```
## # A tibble: 14 x 4  
##   song                artist_name    city      album  
##   <chr>              <chr>      <chr>    <chr>  
## 1 Grievance         Pearl Jam    Seattle, WA Binaural  
## 2 Stupidmop         Pearl Jam    Seattle, WA Vitalogy  
## 3 Present Tense     Pearl Jam    Seattle, WA No Code  
## 4 MFC               Pearl Jam    Seattle, WA Live On Two Legs  
## 5 Lukin             Pearl Jam    Seattle, WA Seattle Washingto~  
## 6 Stuck On Amber    The Boo Radl~ Liverpool,~ Wake Up!  
## 7 It's Lulu         The Boo Radl~ Liverpool,~ Best Of  
## 8 Sparrow           The Boo Radl~ Liverpool,~ Everything's Alri~  
## 9 High as Monkeys   The Boo Radl~ Liverpool,~ Kingsize  
## 10 Butterfly McQueen The Boo Radl~ Liverpool,~ Giant Steps  
## 11 My One and Only Love Carly Simon  New York, ~ Moonlight Serenade  
## 12 It Was So Easy (LP Versio~ Carly Simon  New York, ~ No Secrets  
## 13 I've Got A Crush On You Carly Simon  New York, ~ Clouds In My Coff~  
## 14 "Manha De Carnaval (Theme ~ Carly Simon  New York, ~ Into White
```

1. We really care about the songs in `time`. But, which of those songs do we know its corresponding album?

```
(time %>%  
  inner_join(album, by = "song"))
```

```
## # A tibble: 13 x 6  
##   song                artist_name.x    year artist_name.y city      album  
##   <chr>              <chr>      <dbl> <chr>      <chr>    <chr>  
## 1 Grievance         Pearl Jam    2000 Pearl Jam    Seattle~ Binaural  
## 2 Stupidmop         Pearl Jam    1994 Pearl Jam    Seattle~ Vitalogy  
## 3 Present Tense     Pearl Jam    1996 Pearl Jam    Seattle~ No Code
```

```
## 4 MFC Pearl Jam 1998 Pearl Jam Seattle~ Live On Tw~
## 5 Lukin Pearl Jam 1996 Pearl Jam Seattle~ Seattle Wa~
## 6 It's Lulu The Boo Radle~ 1995 The Boo Radl~ Liverpo~ Best Of
## 7 Sparrow The Boo Radle~ 1992 The Boo Radl~ Liverpo~ Everything~
## 8 High as Monkeys The Boo Radle~ 1998 The Boo Radl~ Liverpo~ Kingsize
## 9 Butterfly McQue~ The Boo Radle~ 1993 The Boo Radl~ Liverpo~ Giant Steps
## 10 My One and Only~ Carly Simon 2005 Carly Simon New Yor~ Moonlight ~
## 11 It Was So Easy ~ Carly Simon 1972 Carly Simon New Yor~ No Secrets
## 12 I've Got A Crus~ Carly Simon 1994 Carly Simon New Yor~ Clouds In ~
## 13 "Manha De Carna~ Carly Simon 2007 Carly Simon New Yor~ Into White
```

2. Go ahead and add the corresponding albums to the time tibble, being sure to preserve rows even if album info is not readily available.

```
(x <- time %>%
  left_join(album, by = "song"))
```

```
## # A tibble: 22 x 6
##   song      artist_name.x year artist_name.y city      album
##   <chr>      <chr>      <dbl> <chr>      <chr>      <chr>
## 1 Corduroy   Pearl Jam    1994 <NA>      <NA>      <NA>
## 2 Grievance  Pearl Jam    2000 Pearl Jam  Seattle~ Binaural
## 3 Stupidmop  Pearl Jam    1994 Pearl Jam  Seattle~ Vitalogy
## 4 Present Tense Pearl Jam    1996 Pearl Jam  Seattle~ No Code
## 5 MFC        Pearl Jam    1998 Pearl Jam  Seattle~ Live On Two ~
## 6 Lukin      Pearl Jam    1996 Pearl Jam  Seattle~ Seattle Wash~
## 7 It's Lulu  The Boo Radle~ 1995 The Boo Radle~ Liverpo~ Best Of
## 8 Sparrow    The Boo Radle~ 1992 The Boo Radle~ Liverpo~ Everything's~
## 9 Martin_ Doom~ The Boo Radle~ 1995 <NA>      <NA>      <NA>
## 10 Leaves And S~ The Boo Radle~ 1993 <NA>      <NA>      <NA>
## # ... with 12 more rows
```

3. Which songs do we have “year”, but not album info?

```
(x %>%
  filter(album == "NA"))
```

```
## # A tibble: 0 x 6
## # ... with 6 variables: song <chr>, artist_name.x <chr>, year <dbl>,
## #   artist_name.y <chr>, city <chr>, album <chr>
```

4. Which artists are in time, but not in album?

```
time %>%
  anti_join(album, by = "artist_name")
```

```
## # A tibble: 5 x 3
##   song      artist_name year
##   <chr>      <chr>      <dbl>
## 1 Mine Again Mariah Carey 2005
## 2 Don't Forget About Us Mariah Carey 2005
## 3 Babydoll   Mariah Carey 1997
## 4 Don't Forget About Us Mariah Carey 2005
## 5 Vision Of Love Mariah Carey 1990
```

5. You’ve come across these two tibbles, and just wish all the info was available in one tibble. What would you do?

```
time %>%
  full_join(album, by = "song")
```

```
## # A tibble: 23 x 6
##   song          artist_name.x year artist_name.y city      album
##   <chr>          <chr>      <dbl> <chr>      <chr>      <chr>
## 1 Corduroy      Pearl Jam    1994 <NA>      <NA>      <NA>
## 2 Grievance     Pearl Jam    2000 Pearl Jam  Seattle~ Binaural
## 3 Stupidmop     Pearl Jam    1994 Pearl Jam  Seattle~ Vitalogy
## 4 Present Tense Pearl Jam    1996 Pearl Jam  Seattle~ No Code
## 5 MFC           Pearl Jam    1998 Pearl Jam  Seattle~ Live On Two ~
## 6 Lukin         Pearl Jam    1996 Pearl Jam  Seattle~ Seattle Wash~
## 7 It's Lulu     The Boo Radle~ 1995 The Boo Radle~ Liverpo~ Best Of
## 8 Sparrow       The Boo Radle~ 1992 The Boo Radle~ Liverpo~ Everything's~
## 9 Martin_ Doom~ The Boo Radle~ 1995 <NA>      <NA>      <NA>
## 10 Leaves And S~ The Boo Radle~ 1993 <NA>      <NA>      <NA>
## # ... with 13 more rows
```

Exercise 2: LOTR

Load in the three Lord of the Rings tibbles that we saw last time:

```
fell <- read_csv("https://raw.githubusercontent.com/jennybc/lotr-tidy/master/data/The_Fellowship_Of_The
```

```
## Parsed with column specification:
## cols(
##   Film = col_character(),
##   Race = col_character(),
##   Female = col_double(),
##   Male = col_double()
## )
```

```
ttow <- read_csv("https://raw.githubusercontent.com/jennybc/lotr-tidy/master/data/The_Two_Towers.csv")
```

```
## Parsed with column specification:
## cols(
##   Film = col_character(),
##   Race = col_character(),
##   Female = col_double(),
##   Male = col_double()
## )
```

```
retk <- read_csv("https://raw.githubusercontent.com/jennybc/lotr-tidy/master/data/The_Return_Of_The_King
```

```
## Parsed with column specification:
## cols(
##   Film = col_character(),
##   Race = col_character(),
##   Female = col_double(),
##   Male = col_double()
## )
```

1. Combine these into a single tibble.

```
full_join(fell, ttow) %>%
  full_join(retk)
```

```
## Joining, by = c("Film", "Race", "Female", "Male")
## Joining, by = c("Film", "Race", "Female", "Male")

## # A tibble: 9 x 4
##   Film                                Race  Female  Male
##   <chr>                             <chr>   <dbl> <dbl>
## 1 The Fellowship Of The Ring Elf      1229   971
## 2 The Fellowship Of The Ring Hobbit    14  3644
## 3 The Fellowship Of The Ring Man       0  1995
## 4 The Two Towers                  Elf     331   513
## 5 The Two Towers                  Hobbit   0  2463
## 6 The Two Towers                  Man     401  3589
## 7 The Return Of The King          Elf     183   510
## 8 The Return Of The King          Hobbit   2  2673
## 9 The Return Of The King          Man     268  2459
```

2. Which races are present in “The Fellowship of the Ring” (`fell`), but not in any of the other ones?

```
fell %>%
  anti_join(ttow, by = "Race") %>%
  anti_join(retk, by = "Race")
```

```
## # A tibble: 0 x 4
## # ... with 4 variables: Film <chr>, Race <chr>, Female <dbl>, Male <dbl>
```

Exercise 3: Set Operations

Let’s use three set functions: `intersect`, `union` and `setdiff`. We’ll work with two toy tibbles named `y` and `z`, similar to Data Wrangling Cheatsheet

```
(y <- tibble(x1 = LETTERS[1:3], x2 = 1:3))
```

```
## # A tibble: 3 x 2
##   x1      x2
##   <chr> <int>
## 1 A         1
## 2 B         2
## 3 C         3
```

```
(z <- tibble(x1 = c("B", "C", "D"), x2 = 2:4))
```

```
## # A tibble: 3 x 2
##   x1      x2
##   <chr> <int>
## 1 B         2
## 2 C         3
## 3 D         4
```

1. Rows that appear in both `y` and `z`

```
union(y, z)
```

```
## # A tibble: 4 x 2
##   x1      x2
##   <chr> <int>
## 1 A         1
## 2 B         2
## 3 C         3
```

```
## 4 D          4
```

2. You collected the data in `y` on Day 1, and `z` in Day 2. Make a data set to reflect that.

```
tibble(  
  mutate(y, day = "Day 1"),  
  mutate(z, day = "Day 2")  
)
```

```
## # A tibble: 3 x 2  
##   `mutate(y, day = "Day 1~` $x2 $day `mutate(z, day = "Day 2~` $x2 $day  
##   <chr>                <int> <chr> <chr>                <int> <chr>  
## 1 A                      1 Day 1 B                      2 Day 2  
## 2 B                      2 Day 1 C                      3 Day 2  
## 3 C                      3 Day 1 D                      4 Day 2
```

3. The rows contained in `z` are bad! Remove those rows from `y`.

```
anti_join(y, z)
```

```
## Joining, by = c("x1", "x2")
```

```
## # A tibble: 1 x 2  
##   x1      x2  
##   <chr> <int>  
## 1 A          1
```